

Youssef Mahmoud 905854027

Alex Hong 905857714

Gedian Wang 705638831

Zachary DeBar 705867064

```

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import PolynomialFeatures
import statsmodels.api as sm
import statsmodels.formula.api as smf
from sklearn.metrics import confusion_matrix, roc_curve, auc, classification_report
from sklearn import preprocessing
from patsy import dmatrix
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns #visualization library
from sklearn.linear_model import LogisticRegression #problem will be solved
from sklearn.metrics import accuracy_score
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis # LDA
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis # QDA
from sklearn.neighbors import KNeighborsClassifier #(KNN)
from sklearn.metrics import confusion_matrix, classification_report, precision_recall_fscore_support
import statsmodels.api as sm #to compute p-values
from patsy import dmatrices
import sklearn.linear_model as skl_lm
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split, LeaveOneOut, KFold, cross_val_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn import tree
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier, export_graphviz
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import confusion_matrix, mean_squared_error
from BorutaShap import BorutaShap
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC, LinearSVC
import matplotlib.pyplot as plt
from matplotlib.font_manager import FontProperties
from matplotlib.patches import Ellipse, Polygon
import matplotlib.gridspec as gridspec
import matplotlib.colors
from pylab import rcParams
from matplotlib.font_manager import FontProperties
from mpl_toolkits.axes_grid1.inset_locator import inset_axes
plt.style.use('seaborn-whitegrid')
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (17, 6)
mpl.rcParams['axes.labelsize'] = 14
mpl.rcParams['xtick.labelsize'] = 12
mpl.rcParams['ytick.labelsize'] = 12
mpl.rcParams['text.color'] = 'k'
%matplotlib inline

import warnings

```

```
warnings.filterwarnings("ignore")
```

```
In [4]: from google.colab import drive
drive.mount('/content/gdrive')
df=pd.read_csv('/content/gdrive/Shareddrives/ECON412/finalproject/smoking.c
df.head()
```

Mounted at /content/gdrive

Out[4]:

	ID	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)	hearing(left)	he
0	0	F	40	155	60	81.3	1.2	1.0	1.0	
1	1	F	40	160	60	81.0	0.8	0.6	1.0	
2	2	M	55	170	60	80.0	0.8	0.8	1.0	
3	3	M	40	165	70	88.0	1.5	1.5	1.0	
4	4	F	40	155	60	86.0	1.0	1.0	1.0	

5 rows × 27 columns

We will now rename and map some of our predictors.

```
In [5]: df1 = df.drop(columns='oral',axis=1)
df1['gender'] = df1['gender'].map({'F':0, 'M':1})
df1['tartar'] = df1['tartar'].map({'N':0, 'Y':1})
df1.rename(columns = {'fasting blood sugar':'fasting_blood_sugar'}, inplace=True)
df1.rename(columns = {'Urine protein':'Urine_protein'}, inplace=True)
df1.rename(columns = {'serum creatinine':'serum_creatinine'}, inplace=True)
df1.rename(columns = {'dental caries':'dental_caries'}, inplace=True)
df1.rename(columns = {'height(cm)':'height'}, inplace=True)
df1.rename(columns = {'weight(kg)':'weight'}, inplace=True)
df1.rename(columns = {'waist(cm)':'waist'}, inplace=True)
df1.rename(columns = {'eyesight(left)':'eyesight_left'}, inplace=True)
df1.rename(columns = {'eyesight(right)':'eyesight_right'}, inplace=True)
df1.rename(columns = {'hearing(left)':'hearing_left'}, inplace=True)
df1.rename(columns = {'hearing(right)':'hearing_right'}, inplace=True)
```

In []: df1

Out[18]:

	ID	gender	age	height	weight	waist	eyesight_left	eyesight_right	hearing_left	hearing_right
0	0	0	40	155	60	81.3	1.2	1.0	1.0	
1	1	0	40	160	60	81.0	0.8	0.6	1.0	
2	2	1	55	170	60	80.0	0.8	0.8	1.0	
3	3	1	40	165	70	88.0	1.5	1.5	1.0	
4	4	0	40	155	60	86.0	1.0	1.0	1.0	
...
55687	55676	0	40	170	65	75.0	0.9	0.9	1.0	
55688	55681	0	45	160	50	70.0	1.2	1.2	1.0	
55689	55683	0	55	160	50	68.5	1.0	1.2	1.0	
55690	55684	1	60	165	60	78.0	0.8	1.0	1.0	
55691	55691	1	55	160	65	85.0	0.9	0.7	1.0	

55692 rows x 26 columns

```
In [ ]: df1.corr()
```

```
Out[19]:
```

	ID	gender	age	height	weight	waist	eyesight_left	
ID	1.000000	0.008657	-0.000825	0.006306	0.004814	0.005384	0.009616	
gender	0.008657	1.000000	-0.290095	0.741556	0.574956	0.419568	0.127424	
age	-0.000825	-0.290095	1.000000	-0.479528	-0.324706	-0.026297	-0.195472	
height	0.006306	0.741556	-0.479528	1.000000	0.675656	0.378902	0.151133	
weight	0.004814	0.574956	-0.324706	0.675656	1.000000	0.822842	0.108433	
waist	0.005384	0.419568	-0.026297	0.378902	0.822842	1.000000	0.027458	
eyesight_left	0.009616	0.127424	-0.195472	0.151133	0.108433	0.027458	1.000000	
eyesight_right	0.003088	0.125680	-0.192723	0.155665	0.113155	0.037996	0.354574	
hearing_left	0.002676	-0.009407	0.203993	-0.078663	-0.050094	0.023790	-0.046571	
hearing_right	-0.004959	-0.011579	0.208722	-0.078323	-0.052836	0.019286	-0.048788	
systolic	0.002489	0.167289	0.134023	0.080585	0.266131	0.316922	-0.019330	
relaxation	0.004649	0.177891	0.050745	0.113193	0.271634	0.292627	0.005199	
fasting_blood_sugar	0.001493	0.098117	0.182351	0.019619	0.136237	0.211132	-0.041851	
Cholesterol	-0.001092	-0.085270	0.055557	-0.082161	0.026403	0.065467	-0.004985	
triglyceride	0.002314	0.241520	0.015102	0.156693	0.324429	0.361922	0.019717	
HDL	-0.005464	-0.306728	0.007047	-0.213284	-0.358868	-0.376203	-0.015296	
LDL	0.001429	-0.042525	0.043007	-0.048419	0.040560	0.072817	-0.007257	
hemoglobin	0.006464	0.702214	-0.263078	0.539367	0.492970	0.387066	0.095234	
Urine_protein	0.000382	0.015907	0.029625	0.005128	0.032566	0.045492	-0.002752	
serum_creatinine	0.003830	0.507249	-0.106118	0.383883	0.324808	0.235024	0.071410	
AST	-0.001865	0.095718	0.032576	0.041737	0.120130	0.142690	-0.007966	
ALT	-0.002803	0.167903	-0.063937	0.126511	0.250634	0.252478	0.019326	
Gtp	0.000823	0.237270	0.013031	0.139720	0.209625	0.243141	0.003850	
dental_caries	0.000641	0.084408	-0.114984	0.079331	0.073536	0.044203	0.003684	
tartar	0.002474	0.055473	-0.081796	0.055513	0.059921	0.046197	0.012532	
smoking	0.011476	0.510340	-0.162557	0.396675	0.302780	0.226259	0.061204	

26 rows × 26 columns

```
In [ ]: , X = dmatrices('smoking ~ age + gender + height + weight + waist + eyesight
logit = sm.Logit(Y, X)
results_logit = logit.fit()
print(results_logit.summary())
```

Optimization terminated successfully.

Current function value: 0.472327

Iterations 7

Logit Regression Results

```
=====
=====
Dep. Variable:          smoking    No. Observations:
55692
Model:                Logit      Df Residuals:
55667
Method:              MLE        Df Model:
24
Date:                Sun, 29 May 2022    Pseudo R-squ.:
0.2816
Time:                14:29:47    Log-Likelihood:          -2
6305.
converged:            True    LL-Null:          -3
6617.
Covariance Type:      nonrobust    LLR p-value:
0.000
=====
=====
```

		coef	std err	z	P> z	[0.0
25	0.975]					

Intercept		-6.7684	0.430	-15.735	0.000	-7.6
11	-5.925					
age		-0.0002	0.001	-0.158	0.875	-0.0
02	0.002					
gender		2.9168	0.051	56.774	0.000	2.8
16	3.017					
height		0.0225	0.002	10.117	0.000	0.0
18	0.027					
weight		-0.0103	0.002	-4.776	0.000	-0.0
15	-0.006					
waist		-0.0012	0.003	-0.465	0.642	-0.0
06	0.004					
eyesight_left		-0.0196	0.023	-0.870	0.384	-0.0
64	0.025					
eyesight_right		-0.0103	0.023	-0.444	0.657	-0.0
56	0.035					
hearing_left		-0.2188	0.081	-2.688	0.007	-0.3
78	-0.059					
hearing_right		0.0213	0.080	0.264	0.792	-0.1
36	0.179					
systolic		-0.0144	0.001	-11.401	0.000	-0.0
17	-0.012					
relaxation		0.0097	0.002	5.604	0.000	0.0
06	0.013					
fasting_blood_sugar		0.0035	0.001	6.367	0.000	0.0
02	0.005					

Cholesterol	-0.0024	0.001	-4.619	0.000	-0.0
03 -0.001					
triglyceride	0.0047	0.000	23.358	0.000	0.0
04 0.005					
HDL	0.0021	0.001	2.105	0.035	0.0
00 0.004					
LDL	-0.0001	0.000	-0.350	0.726	-0.0
01 0.001					
hemoglobin	0.1390	0.011	12.745	0.000	0.1
18 0.160					
Urine_protein	0.0117	0.027	0.439	0.660	-0.0
41 0.064					
serum_creatinine	-0.8690	0.067	-12.949	0.000	-1.0
00 -0.737					
AST	-0.0011	0.001	-1.050	0.294	-0.0
03 0.001					
ALT	-0.0054	0.001	-6.835	0.000	-0.0
07 -0.004					
Gtp	0.0073	0.000	22.203	0.000	0.0
07 0.008					
dental_caries	0.3163	0.026	12.285	0.000	0.2
66 0.367					
tartar	0.3376	0.022	15.472	0.000	0.2
95 0.380					
=====					
=====					

Based on the P-values, the following variables are significant: gender, height, weight, relaxation, systolic, fasting_blood_sugar, Cholesterol, triglyceride, HDL, hemoglobin, serum_creatinine, ALT, Gtp, dental_caries & tartar. Now we will run another logit using only the significant predictors.

```
In [ ]: Y, X = dmatrices('smoking ~ gender + height + weight + systolic + relaxatio
logit1 = sm.Logit(Y, X)
results_logit1 = logit1.fit()
print(results_logit1.summary())
```

Optimization terminated successfully.

Current function value: 0.472436

Iterations 7

Logit Regression Results

```
=====
=====
Dep. Variable:          smoking    No. Observations:
55692
Model:                Logit      Df Residuals:
55676
Method:              MLE        Df Model:
15
Date:                Sun, 29 May 2022    Pseudo R-squ.:
0.2815
Time:                14:29:48    Log-Likelihood:          -2
6311.
converged:            True    LL-Null:          -3
6617.
Covariance Type:      nonrobust    LLR p-value:
0.000
=====
=====
```

		coef	std err	z	P> z	[0.0
25	0.975]					

Intercept		-7.1971	0.348	-20.671	0.000	-7.8
79	-6.515					
gender		2.9039	0.051	57.059	0.000	2.8
04	3.004					
height		0.0232	0.002	11.450	0.000	0.0
19	0.027					
weight		-0.0109	0.001	-8.584	0.000	-0.0
13	-0.008					
systolic		-0.0146	0.001	-11.700	0.000	-0.0
17	-0.012					
relaxation		0.0098	0.002	5.702	0.000	0.0
06	0.013					
fasting_blood_sugar		0.0035	0.001	6.467	0.000	0.0
02	0.005					
Cholesterol		-0.0026	0.000	-7.799	0.000	-0.0
03	-0.002					
triglyceride		0.0048	0.000	25.609	0.000	0.0
04	0.005					
HDL		0.0023	0.001	2.507	0.012	0.0
01	0.004					
hemoglobin		0.1408	0.011	13.130	0.000	0.1
20	0.162					
serum_creatinine		-0.8733	0.067	-13.094	0.000	-1.0
04	-0.743					
ALT		-0.0060	0.001	-10.572	0.000	-0.0
07	-0.005					

Gtp		0.0073	0.000	22.287	0.000	0.0
07	0.008					
dental_caries		0.3184	0.026	12.430	0.000	0.2
68	0.369					
tartar		0.3398	0.022	15.600	0.000	0.2
97	0.382					

=====

=====

```
In [ ]: X = df1.drop(['smoking'], axis = 1)
Y = df1['smoking']
X_train1, X_test1, Y_train1, Y_test1 = train_test_split(X, Y, test_size=0.0


Feature_Selector = BorutaShap(importance_measure='shap', classification=True
Feature_Selector.fit(X=X_test1, y=Y_test1, n_trials=20, random_state=0)
Feature_Selector.plot(which_features='all')
```

0% | 0/20 [00:00<?, ?it/s]

6 attributes confirmed important: ['gender', 'height', 'triglyceride', 'HDL', 'Gtp', 'hemoglobin']

19 attributes confirmed unimportant: ['eyesight_right', 'ALT', 'weight', 'relaxation', 'hearing_right', 'ID', 'tartar', 'hearing_left', 'systolic', 'Cholesterol', 'dental_caries', 'Urine_protein', 'eyesight_left', 'serum_creatinine', 'age', 'fasting_blood_sugar', 'LDL', 'AST', 'waist']

0 tentative attributes remains: []



All of our chosen variables are statistically significant, now we will proceed with our analysis.

```
In [ ]: X1 = df1[['gender', 'height', 'weight', 'relaxation', 'systolic', 'fasting', 'smoking']]
        Y1 = df1.smoking
```

```
In [ ]: X1_train, X1_test, Y1_train, Y1_test = sklearn.model_selection.train_test_split(X1, Y1, test_size=0.2, random_state=42)
```

```
In [ ]: lr = LogisticRegression()
```

```
In [ ]: logit2 = lr.fit(X1_train, Y1_train)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, lr.predict(X1_test))
        print(conf_mat)
```

```
lr.score(X1_test, Y1_test)
print('Test Accuracy =', lr.score(X1_test, Y1_test))
```

```
[[7051 1814]
 [2072 2986]]
Test Accuracy = 0.720893485599368
```

LDA

```
In [ ]: lda = LinearDiscriminantAnalysis()
        lda.fit(X1_train, Y1_train)
        LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None,
                                     solver='svd', store_covariance=False, tol=0.0001)
```

```
Out[27]: LinearDiscriminantAnalysis()
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, lda.predict(X1_test))
        print(conf_mat)
        lda.score(X1_test, Y1_test)
        print('Test Accuracy =', lda.score(X1_test, Y1_test))
```

```
[[6515 2350]
 [1203 3855]]
Test Accuracy = 0.7448107448107448
```

LDA performed better than the Logit Model.

QDA

```
In [ ]: qda = QuadraticDiscriminantAnalysis()
        qda.fit(X1_train, Y1_train)
        QuadraticDiscriminantAnalysis(priors=None, reg_param=0.0,
                                       store_covariance=False, tol=0.0001)
```

```
Out[29]: QuadraticDiscriminantAnalysis()
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, qda.predict(X1_test))
print(conf_mat)
qda.score(X1_test, Y1_test)
print('Test Accuracy =', qda.score(X1_test, Y1_test))

[[6528 2337]
 [1374 3684]]
Test Accuracy = 0.7334626158155569
```

LDA performed better than QDA, implying that our classes might not require a non-linear classifier.

Naive Bayes

```
In [ ]: from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X1_train, Y1_train)
```

```
Out[31]: GaussianNB()
```

```
In [ ]: Y1_pred = classifier.predict(X1_test)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, Y1_pred)
print(conf_mat)
accuracy_score(Y1_test, Y1_pred)
print('Test Accuracy =', accuracy_score(Y1_test, Y1_pred))

[[5582 3283]
 [ 804 4254]]
Test Accuracy = 0.7064569417510594
```

Naive Bayes performs poorly when compared to LDA, QDA and Logit.

KNN

```
In [ ]: nbrs1 = KNeighborsClassifier(n_neighbors=1)
nbrs1.fit(X1_train, Y1_train)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                    weights='uniform')
```

```
Out[7]: KNeighborsClassifier(n_jobs=1, n_neighbors=1)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, nbrs1.predict(X1_test))
print(conf_mat)
nbrs1.score(X1_test, Y1_test)
print('Test Accuracy =', nbrs1.score(X1_test, Y1_test))

[[7192 1673]
 [1725 3333]]
Test Accuracy = 0.7559434030022265
```

```
In [ ]: nbrs2 = KNeighborsClassifier(n_neighbors=2)
nbrs2.fit(X1_train,Y1_train)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=2, p=2,
                    weights='uniform')
```

```
Out[36]: KNeighborsClassifier(n_jobs=1, n_neighbors=2)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, nbrs2.predict(X1_test))
print(conf_mat)
nbrs2.score(X1_test, Y1_test)
print('Test Accuracy =', nbrs2.score(X1_test, Y1_test))

[[7909  956]
 [3043 2015]]
Test Accuracy = 0.7127774186597716
```

```
In [ ]: nbrs3 = KNeighborsClassifier(n_neighbors=3)
nbrs3.fit(X1_train,Y1_train)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=3, p=2,
                    weights='uniform')
```

```
Out[38]: KNeighborsClassifier(n_jobs=1, n_neighbors=3)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, nbrs3.predict(X1_test))
print(conf_mat)
nbrs3.score(X1_test, Y1_test)
print('Test Accuracy =', nbrs3.score(X1_test, Y1_test))

[[6902 1963]
 [2105 2953]]
Test Accuracy = 0.7078215901745314
```

```
In [ ]: nbrs4 = KNeighborsClassifier(n_neighbors=4)
nbrs4.fit(X1_train,Y1_train)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=4, p=2,
                    weights='uniform')
```

```
Out[40]: KNeighborsClassifier(n_jobs=1, n_neighbors=4)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, nbrs4.predict(X1_test))
print(conf_mat)
nbrs4.score(X1_test, Y1_test)
print('Test Accuracy =', nbrs4.score(X1_test, Y1_test))

[[7645 1220]
 [2841 2217]]
Test Accuracy = 0.708324355383179
```

```
In [ ]: nbrs5 = KNeighborsClassifier(n_neighbors=5)
nbrs5.fit(X1_train,Y1_train)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                    weights='uniform')
```

```
Out[42]: KNeighborsClassifier(n_jobs=1)
```

```
In [ ]: conf_mat = confusion_matrix(Y1_test, nbrs5.predict(X1_test))
print(conf_mat)
nbrs5.score(X1_test, Y1_test)
print('Test Accuracy =', nbrs5.score(X1_test, Y1_test))
```

```
[[6933 1932]
 [2134 2924]]
Test Accuracy = 0.7079652373770021
```

KNN with 1 neighbor performed better when compared to other KNNs and all the models ran so far.


Decision Tree Classifier

```
In [ ]: tree_depth = []

for i in range(1,21):
    cv_tree = DecisionTreeClassifier(max_depth=i)
    scores = cross_val_score(estimator=cv_tree, X=X1_train, y=Y1_train, cv=5)
    tree_depth.append(scores.mean())

plt.plot(range(1,21), tree_depth)
```

```
Out[9]: [<matplotlib.lines.Line2D at 0x7fb2709eb280>]
```



According to the cross validation plot, the tree depth that produces the lowest training MSE is around 20.

```
In [ ]: clf = DecisionTreeClassifier(max_depth=20)
        clf.fit(X1_train, Y1_train)
        print("Training Accuracy =", clf.score(X1_train, Y1_train))
```

Training Accuracy = 0.9700256170844406

```
In [ ]: pred1 = clf.predict(X1_test)
        cm = pd.DataFrame(confusion_matrix(Y1_test, pred1).T, index=['No', 'Yes'],
                           print(cm)
                           print('Test Accuracy =', (10707/13923))
```

	No	Yes
No	7203	1555
Yes	1662	3503


Test Accuracy = 0.7690152984270632

```
In [ ]: tree_depth = []

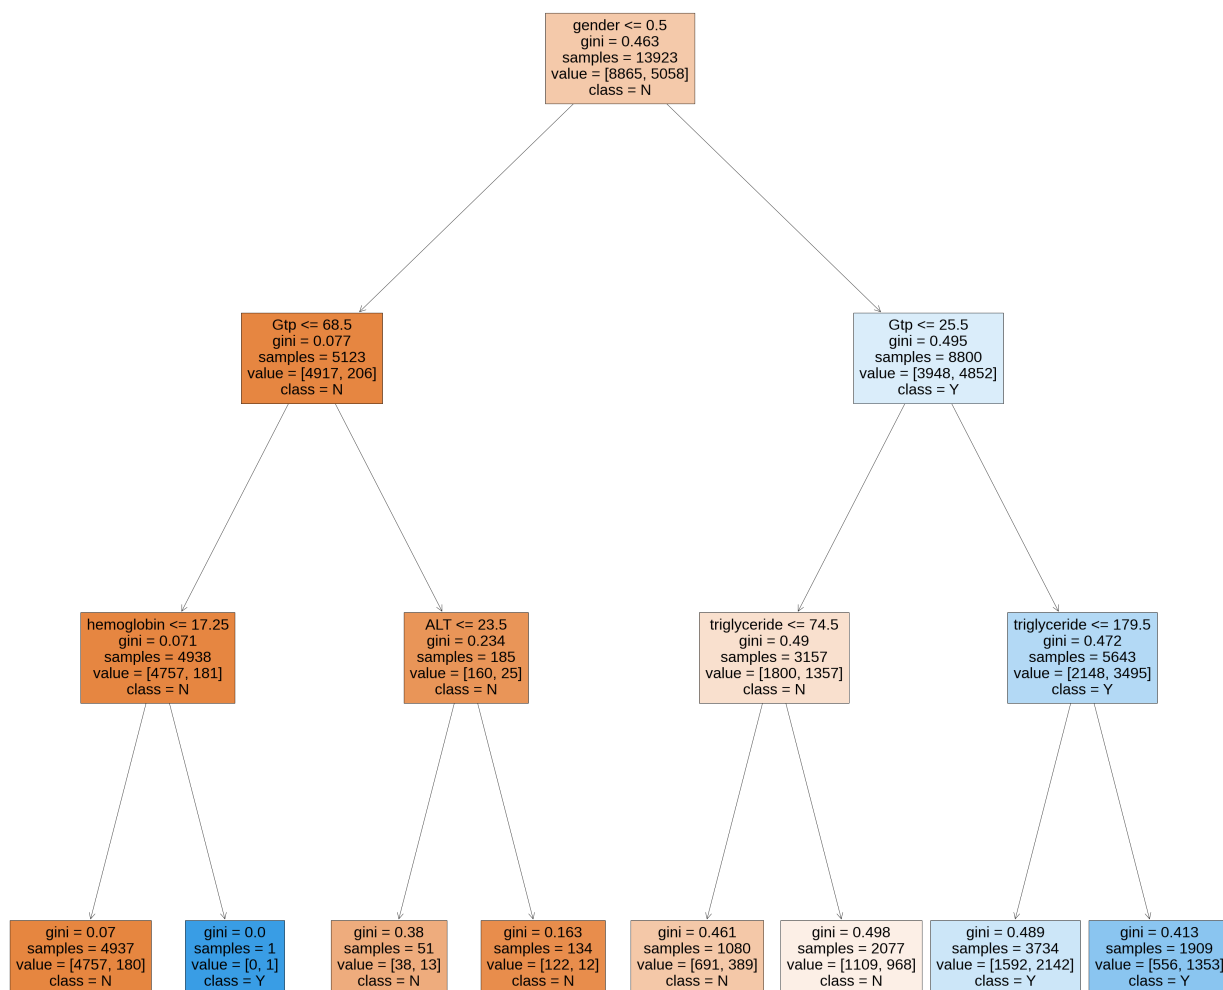
        for i in range(1,21):
            cv_tree = DecisionTreeClassifier(max_depth=i)
            scores = cross_val_score(estimator=cv_tree, X=X1_test, y=Y1_test, cv=1)
            tree_depth.append(scores.mean())

        plt.plot(range(1,21), tree_depth)
```

Out[12]: [<matplotlib.lines.Line2D at 0x7fb26ce850a0>]



```
In [ ]: fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (60,60))
tree.plot_tree(clf1,feature_names = X1_test.columns,filled = True, class_na
```



So far, Decision Tree Classifier performed the best with a test accuracy of approximately 77%

SVC

```
In [ ]: svc1 = SVC(C = 0.01)
Fit1 = svc1.fit(X1_train, Y1_train)
```

```
In [ ]: pred2 = svc1.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred2).T, index=['No', 'Yes'],
print(cm)
print('Test Accuracy =',(9606/13923))
```

	No	Yes
No	8047	3499
Yes	818	1559

Test Accuracy = 0.6899375134669252

```
In [ ]: svc2 = SVC(C = 0.1)
Fit2 = svc2.fit(X1_train, Y1_train)
```

```
In [ ]: pred3 = svc2.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred3).T, index=['No', 'Yes'],
print(cm)
print('Test Accuracy =',(10005/13923))
```

	No	Yes
No	7687	2740
Yes	1178	2318

Test Accuracy = 0.7185951303598362

```
In [ ]: svc3 = SVC(C = 1)
Fit3 = svc3.fit(X1_train, Y1_train)
```

```
In [ ]: pred4 = svc3.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred4).T, index=['No', 'Yes'],
print(cm)
print('Test Accuracy =',(10301/13923))
```

	No	Yes
No	7454	2211
Yes	1411	2847

Test Accuracy = 0.7398549163255046

```
In [ ]: svc4 = SVC(C = 5)
Fit4 = svc4.fit(X1_train, Y1_train)
```

```
In [ ]: pred5 = svc4.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred5).T, index=['No', 'Yes'],
print(cm)
print('Test Accuracy =',(10389/13923))
```

	No	Yes
No	7322	1991
Yes	1543	3067

Test Accuracy = 0.7461753932342168


```
In [ ]: svc5 = SVC(C = 10)
Fit5 = svc5.fit(X1_train, Y1_train)
```

```
In [ ]: pred6 = svc5.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred6).T, index=['No', 'Yes'],
print(cm)
print('Test Accuracy =', (10408/13923))
```

	No	Yes
No	7237	1887
Yes	1628	3171

Test Accuracy = 0.7475400416576887

SVC with a linear kernel and C=10 performed better than SVCs with lower values of C, however, it still could not beat the performance of the Decision Tree Classifier.

SVM

```
In [ ]: svm1 = SVC(C=1, kernel='rbf')
Fit6 = svm1.fit(X1_train, Y1_train)
```

```
In [ ]: pred7 = svm1.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred7).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7454	2211
Yes	1411	2847

```
In [ ]: print('Test Accuracy =', ((7454+2847)/13923))
```

Test Accuracy = 0.7398549163255046

```
In [ ]: svm2 = SVC(C=3, kernel='rbf')
Fit7 = svm2.fit(X1_train, Y1_train)
```

```
In [ ]: pred8 = svm2.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred8).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7365	2063
Yes	1500	2995

```
In [ ]: print('Test Accuracy =', ((7365+2995)/13923))
```

Test Accuracy = 0.7440925087983912

```
In [ ]: svm3 = SVC(C=5, kernel='rbf')
Fit8 = svm3.fit(X1_train, Y1_train)
```

```
In [ ]: pred9 = svm3.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred9).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7322	1991
Yes	1543	3067

```
In [ ]: print('Test Accuracy =', ((7322+3067)/13923))
```

Test Accuracy = 0.7461753932342168

```
In [ ]: svm4 = SVC(C=8, kernel='rbf')
Fit9 = svm4.fit(X1_train, Y1_train)
```

```
In [ ]: pred10 = svm4.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred10).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7268	1927
Yes	1597	3131

```
In [ ]: print('Test Accuracy =', ((7268+3131)/13923))
```

Test Accuracy = 0.7468936292465704

```
In [ ]: svm5 = SVC(C=10, kernel='rbf')
Fit10 = svm5.fit(X1_train, Y1_train)
```

```
In [ ]: pred11 = svm5.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred11).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7237	1887
Yes	1628	3171

```
In [ ]: print('Test Accuracy =', ((7237+3171)/13923))
```

Test Accuracy = 0.7475400416576887

```
In [ ]: svm6 = SVC(C=1, kernel='poly')
Fit11 = svm6.fit(X1_train, Y1_train)
```

```
In [ ]: pred12 = svm6.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred12).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7361	2093
Yes	1504	2965

```
In [ ]: print('Test Accuracy =', ((7361+2965)/13923))

Test Accuracy = 0.7416505063563887
```

```
In [ ]: svm10 = SVC(C=3, kernel='poly')
Fit14 = svm10.fit(X1_train, Y1_train)
```

```
In [ ]: pred16 = svm10.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred16).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7264	1945
Yes	1601	3113

```
In [ ]: print('Test Accuracy =', ((7264+3113)/13923))

Test Accuracy = 0.7453135100193924
```

```
In [ ]: svm7 = SVC(C=5, kernel='poly')
Fit12 = svm7.fit(X1_train, Y1_train)
```

```
In [ ]: pred13 = svm7.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred13).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7189	1826
Yes	1676	3232

```
In [ ]: print('Test Accuracy =', ((7189+3232)/13923))

Test Accuracy = 0.7484737484737485
```

```
In [ ]: svm8 = SVC(C=8, kernel='poly')
Fit13 = svm8.fit(X1_train, Y1_train)
```

```
In [ ]: pred14 = svm8.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred14).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7136	1747
Yes	1729	3311

```
In [ ]: print('Test Accuracy =', ((7136+3311)/13923))
```

Test Accuracy = 0.750341162105868

```
In [ ]: svm9 = SVC(C=10, kernel='poly')
Fit13 = svm9.fit(X1_train, Y1_train)
```

```
In [ ]: pred15 = svm9.predict(X1_test)
cm = pd.DataFrame(confusion_matrix(Y1_test, pred15).T, index=['No', 'Yes'],
print(cm)
```

	No	Yes
No	7091	1703
Yes	1774	3355

```
In [ ]: print('Test Accuracy =', ((7091+3355)/13923))
```


Test Accuracy = 0.7502693385046326

SVM with C=8 and a polynomial kernel performed best when compared to other SVMs and SVCs, however, it still couldnt beat the performance of the Decision Tree Classifier.

```
In [ ]: x = preprocessing.scale(X1)
y = np.ravel(Y1)
```

```
In [ ]: fig = plt.figure()

fig.suptitle('decison surface using projected features')
labels = ['SVC', 'SVM poly', 'SVM radial' ]
gs = gridspec.GridSpec(1, 1)
for svm8, lab, grd in zip([Fit13], labels, ([0,0], [1,0], [2,0])):
    svm8.fit(np.stack((x[:,0], x[:,1]), axis=-1), y)
    ax = plt.subplot(gs[grd[0], grd[1]])
    fig = plot_decision_regions(X=np.stack((x[:,0], x[:,1]), axis=-1), y=y,
    plt.title(lab)
plt.show()
```



```

In [ ]: plt.style.use('seaborn-whitegrid')
fig, ax = plt.subplots(1, figsize=(15,6))
# false positive rates and true positive rates
fpr, tpr, _ = roc_curve(Y1_test, clf.predict(X1_test))
fpr1, tpr1, _ = roc_curve(Y1_test, nbrs1.predict(X1_test))
_ = ax.plot(fpr, tpr, lw=2, label='Classification Tree. ROC curve (Area = %
_ = ax.plot(fpr1, tpr1, lw=2, label='KNN. ROC curve (Area = %0.2f)' % auc(f
_ = ax.set_title('Classification Tree and KNN n=1')
_ = ax.plot([0, 1], 'k--', lw=2)
_ = ax.set_xlim([-0.05, 1.0])
_ = ax.set_ylim([0.0, 1.05])
_ = ax.set_xlabel('False Positive Rate (FPR)')
_ = ax.set_ylabel('True Positive Rate (TPR)')
_ = ax.legend(loc='center left', bbox_to_anchor=(.05, -0.3), fontsize = 12)
_ = ax.set_aspect(1)

```

