

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 from statsmodels.tsa.arima_process import ArmaProcess
4 import matplotlib.pyplot as plt
5 import datetime
6 import io
7 import datetime
8 import matplotlib.lines as mlines
9 import statsmodels.formula.api as smf
10 from statsmodels.tsa.arima.model import ARIMA
11 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
12 from statsmodels.tsa.stattools import adfuller
13 from pmdarima.arima import auto_arima
14 from statsmodels.tsa.arima.model import ARIMA
15 import scipy.stats as st
16 import pmdarima as pm
```

# 1

```
In [2]: 1 df = pd.read_csv("desktop/JPYUSD.csv", parse_dates = True, index_col
```

```
In [3]: 1 df
```

Out[3]:

	Japan
1985-01-01	0.003927
1985-02-01	0.003854
1985-03-01	0.003960
1985-04-01	0.003964
1985-05-01	0.003971
...	...
2022-07-01	0.007519
2022-08-01	0.007214
2022-09-01	0.006909
2022-10-01	0.006746
2022-11-01	0.007205

455 rows × 1 columns

```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 455 entries, 1985-01-01 to 2022-11-01
Data columns (total 1 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Japan   455 non-null     float64
dtypes: float64(1)
memory usage: 7.1 KB
```

```
In [5]: 1 df80 = df.iloc[:int(len(df)*0.8)]
```

```
In [6]: 1 df80
```

```
Out[6]:
```

	Japan
1985-01-01	0.003927
1985-02-01	0.003854
1985-03-01	0.003960
1985-04-01	0.003964
1985-05-01	0.003971
...	...
2014-12-01	0.008289
2015-01-01	0.008459
2015-02-01	0.008385
2015-03-01	0.008326
2015-04-01	0.008438

364 rows × 1 columns

```
In [7]: 1 plt.plot(df80)
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x7fc2c4e20d30>]
```



```
In [8]: 1 df80.diff()
```

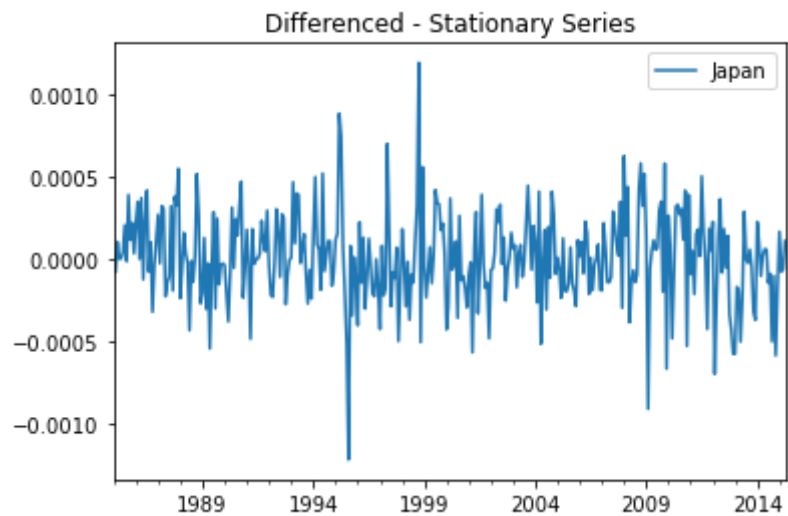
```
Out[8]:
```

Japan	
1985-01-01	NaN
1985-02-01	-0.000073
1985-03-01	0.000107
1985-04-01	0.000004
1985-05-01	0.000006
...	...
2014-12-01	-0.000170
2015-01-01	0.000170
2015-02-01	-0.000074
2015-03-01	-0.000059
2015-04-01	0.000112

364 rows × 1 columns

```
In [9]: 1 df80.diff().plot()  
        2 plt.title("Differenced - Stationary Series")
```

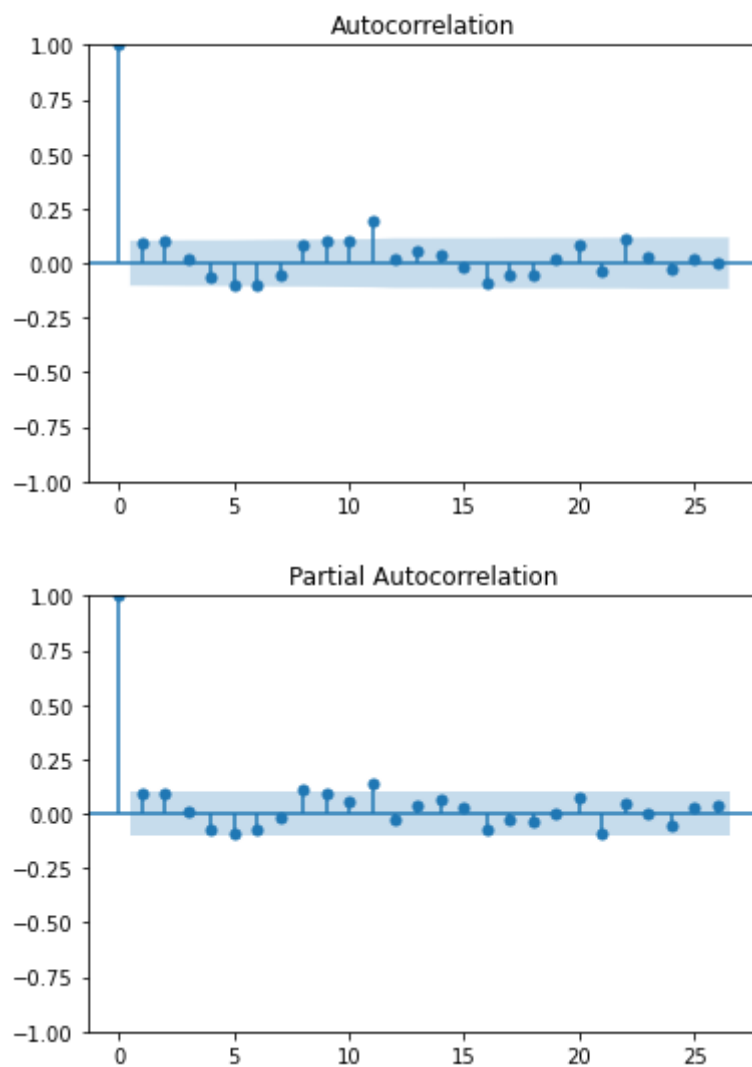
```
Out[9]: Text(0.5, 1.0, 'Differenced - Stationary Series')
```



The differenced training data looks stationary as it is mean converting.

```
In [10]: 1 plot_acf(df80['Japan'].diff().dropna())  
2 plot_pacf(df80['Japan'].diff().dropna(), method='ywm')  
3 plt.plot()
```

Out[10]: []



```
In [ ]: 1
```

```
In [11]: 1 model3 = ARIMA(df80.dropna(), order=(1,1,0)).fit()
          2 model3.summary()
```

```
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/stats
models/tsa/base/tsa_model.py:471: ValueWarning: No frequency informat
ion was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/stats
models/tsa/base/tsa_model.py:471: ValueWarning: No frequency informat
ion was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/stats
models/tsa/base/tsa_model.py:471: ValueWarning: No frequency informat
ion was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/stats
models/base/model.py:604: ConvergenceWarning: Maximum Likelihood opti
mization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to ")
```

Out[11]: SARIMAX Results

<b>Dep. Variable:</b>	Japan	<b>No. Observations:</b>	364
<b>Model:</b>	ARIMA(1, 1, 0)	<b>Log Likelihood</b>	2461.065
<b>Date:</b>	Fri, 27 Jan 2023	<b>AIC</b>	-4918.130
<b>Time:</b>	16:19:55	<b>BIC</b>	-4910.341
<b>Sample:</b>	01-01-1985	<b>HQIC</b>	-4915.034
	- 04-01-2015		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.0991	0.044	2.267	0.023	0.013	0.185
<b>sigma2</b>	7.537e-08	4.21e-09	17.901	0.000	6.71e-08	8.36e-08

<b>Ljung-Box (L1) (Q):</b>	0.04	<b>Jarque-Bera (JB):</b>	41.03
<b>Prob(Q):</b>	0.83	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	1.49	<b>Skew:</b>	-0.09
<b>Prob(H) (two-sided):</b>	0.03	<b>Kurtosis:</b>	4.64

Warnings:

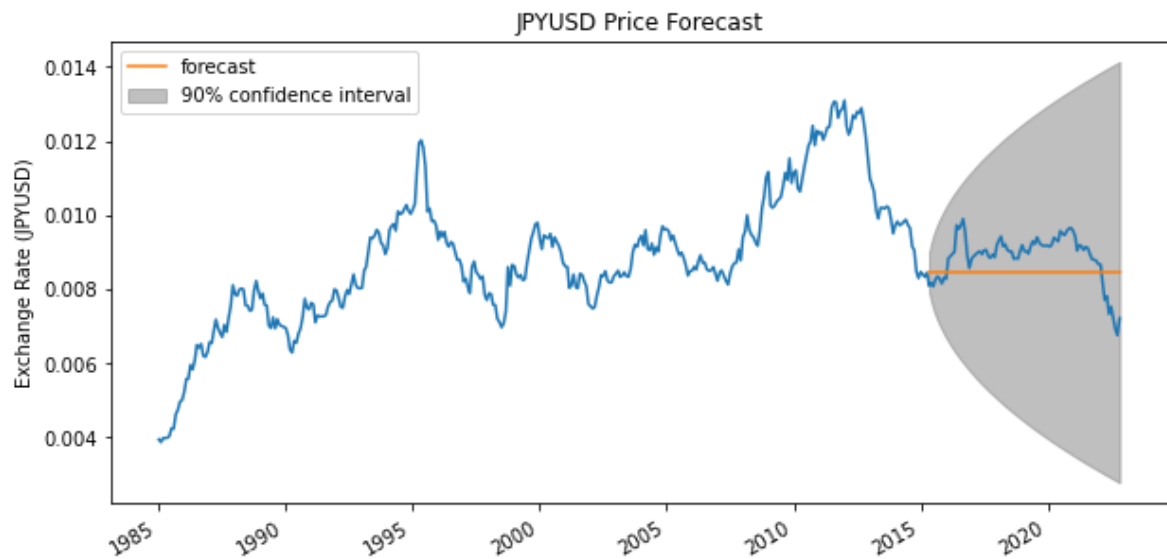
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [12]: 1 forecast2 = model3.get_forecast(12)
          2
          3 # get the 90% confidence interval for the forecast
          4 yhat_conf_int2 = forecast2.conf_int(alpha=0.1)
          5 yhat_conf_int2["mean"] = forecast2.predicted_mean
          6 yhat_conf_int2
```

```
Out[12]:
```

	lower Japan	upper Japan	mean
<b>2015-05-01</b>	0.007998	0.008901	0.008449
<b>2015-06-01</b>	0.007779	0.009121	0.008450
<b>2015-07-01</b>	0.007613	0.009288	0.008450
<b>2015-08-01</b>	0.007475	0.009426	0.008450
<b>2015-09-01</b>	0.007353	0.009548	0.008450
<b>2015-10-01</b>	0.007244	0.009657	0.008450
<b>2015-11-01</b>	0.007144	0.009757	0.008450
<b>2015-12-01</b>	0.007051	0.009850	0.008450
<b>2016-01-01</b>	0.006964	0.009937	0.008450
<b>2016-02-01</b>	0.006882	0.010019	0.008450
<b>2016-03-01</b>	0.006804	0.010097	0.008450
<b>2016-04-01</b>	0.006729	0.010172	0.008450

```
In [13]: 1 from statsmodels.graphics.tsaplots import plot_predict
2
3 # Plot the data and the forecast
4 fig, ax = plt.subplots(figsize = (10, 5))
5 plt.title("JPYUSD Price Forecast")
6 plt.plot(df['Japan'])
7 plt.ylabel("Exchange Rate (JPYUSD)")
8 plot_predict(model3, ax=ax, start = "2015-05-01", end = "2022-11-01")
9 plt.legend()
10 plt.show()
```





```

In [14]: 1 def evaluate_arima_model(X, arima_order):
2         # prepare training dataset
3         train_size = int(len(X) * 0.8)
4         train, test = X[0:train_size], X[train_size:]
5         history = [x for x in train]
6         # make predictions
7         predictions = list()
8         for t in range(len(test)):
9             model = ARIMA(history, order=arima_order)
10            model_fit = model.fit()
11            yhat = model_fit.forecast()[0]
12            predictions.append(yhat)
13            history.append(test[t])
14        # calculate out of sample error
15        test = pd.DataFrame(test)
16        test["predictions"] = predictions
17        return test
18 y = evaluate_arima_model(df['Japan'], (1,1,0))

```

```

/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/st
atsmodels/base/model.py:604: ConvergenceWarning: Maximum Likeliho
d optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/st
atsmodels/base/model.py:604: ConvergenceWarning: Maximum Likeliho
d optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/st
atsmodels/base/model.py:604: ConvergenceWarning: Maximum Likeliho
d optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/st
atsmodels/base/model.py:604: ConvergenceWarning: Maximum Likeliho
d optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
/Users/youssefmahmoud/opt/anaconda3/lib/python3.9/site-packages/st
atsmodels/base/model.py:604: ConvergenceWarning: Maximum Likeliho
d optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "

```

```
In [15]: 1 y
```

```
Out[15]:
```

	Japan	predictions
2015-05-01	0.008081	0.008449
2015-06-01	0.008167	0.008046
2015-07-01	0.008066	0.008175
2015-08-01	0.008251	0.008057
2015-09-01	0.008337	0.008268
...	...	...
2022-07-01	0.007519	0.007272
2022-08-01	0.007214	0.007538
2022-09-01	0.006909	0.007186
2022-10-01	0.006746	0.006880
2022-11-01	0.007205	0.006730

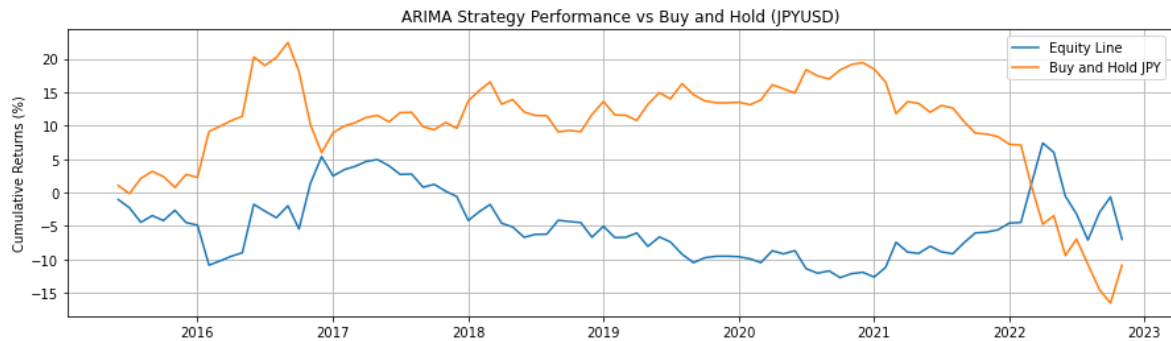
91 rows × 2 columns

```
In [16]: 1 y["Signals"] = np.where(y["predictions"]>y['Japan'].shift(), 1,-1)
2
3
4 y["returns"] = np.log(y["Japan"]/y["Japan"].shift())
5
6
7 y["strategy returns"] = y["Signals"]*y["returns"]
8
9
10 y["Cumulative Returns"] = (np.exp(y["strategy returns"].cumsum())-1
```

```
In [17]: 1 y["Cumulative Returns"]
```

```
Out[17]: 2015-05-01      NaN
2015-06-01    -1.050505
2015-07-01    -2.263728
2015-08-01    -4.447558
2015-09-01    -3.451806
...
2022-07-01    -3.187663
2022-08-01    -7.105975
2022-09-01    -2.997755
2022-10-01    -0.652115
2022-11-01    -6.978640
Name: Cumulative Returns, Length: 91, dtype: float64
```

```
In [18]: 1 plt.figure(figsize = (15, 4))
2 plt.plot(y["Cumulative Returns"])
3 plt.plot((np.exp(y["returns"].cumsum())-1)*100)
4 plt.ylabel("Cumulative Returns (%)")
5 plt.title("ARIMA Strategy Performance vs Buy and Hold (JPYUSD)")
6 plt.legend(["Equity Line", "Buy and Hold JPY"])
7 plt.grid()
```



```
In [19]: 1 y["Cumulative Returns"]/(100+1)
```

```
Out[19]: 2015-05-01      NaN
2015-06-01    -0.010401
2015-07-01    -0.022413
2015-08-01    -0.044035
2015-09-01    -0.034176
...
2022-07-01    -0.031561
2022-08-01    -0.070356
2022-09-01    -0.029681
2022-10-01    -0.006457
2022-11-01    -0.069095
Name: Cumulative Returns, Length: 91, dtype: float64
```

```
In [20]: 1 P = 1000000
2 A = (((y["Cumulative Returns"]/100)+1)*P)[-1]
3 t = (len(y)/12)
4 CCROR = np.log(A/P)/t
5 print(CCROR*100)
```

```
-0.9539477975994095
```

```
In [21]: 1 ((A/P)**(1/t)-1)*100 # Annualized
```

```
Out[21]: -0.9494121496260166
```

## 2

```
In [22]: 1 y['error'] = y['Japan'] - y['predictions']
          2 y
```

Out[22]:

	Japan	predictions	Signals	returns	strategy returns	Cumulative Returns	error
2015-05-01	0.008081	0.008449	-1	NaN	NaN	NaN	-0.000368
2015-06-01	0.008167	0.008046	-1	0.010561	-0.010561	-1.050505	0.000121
2015-07-01	0.008066	0.008175	1	-0.012337	-0.012337	-2.263728	-0.000108
2015-08-01	0.008251	0.008057	-1	0.022598	-0.022598	-4.447558	0.000194
2015-09-01	0.008337	0.008268	1	0.010367	0.010367	-3.451806	0.000068
...	...	...	...	...	...	...	...
2022-07-01	0.007519	0.007272	-1	0.026927	-0.026927	-3.187663	0.000247
2022-08-01	0.007214	0.007538	1	-0.041315	-0.041315	-7.105975	-0.000323
2022-09-01	0.006909	0.007186	-1	-0.043275	0.043275	-2.997755	-0.000277
2022-10-01	0.006746	0.006880	-1	-0.023894	0.023894	-0.652115	-0.000134
2022-11-01	0.007205	0.006730	-1	0.065799	-0.065799	-6.978640	0.000475

91 rows × 7 columns

```
In [23]: 1 s_current = np.log(y["Japan"]).reset_index(drop=True)
          2 s_current = s_current.rename('s_current')
          3 s_current
```

Out[23]:

0	-4.818263
1	-4.807703
2	-4.820040
3	-4.797442
4	-4.787075
...	
86	-4.890349
87	-4.931664
88	-4.974939
89	-4.998833
90	-4.933034

Name: s\_current, Length: 91, dtype: float64

```
In [24]: 1 df3 = pd.DataFrame(s_current, columns = ['s_current'])
          2
```

```
In [25]: 1 df3
```

```
Out[25]:
```

	s_current
0	-4.818263
1	-4.807703
2	-4.820040
3	-4.797442
4	-4.787075
...	...
86	-4.890349
87	-4.931664
88	-4.974939
89	-4.998833
90	-4.933034

```
In [26]: 1 s_future = s_current.shift(1)
```

```
In [27]: 1 s_future.fillna(0)
```

```
Out[27]:
```

0	0.000000
1	-4.818263
2	-4.807703
3	-4.820040
4	-4.797442
...	...
86	-4.917277
87	-4.890349
88	-4.931664
89	-4.974939
90	-4.998833

Name: s\_current, Length: 91, dtype: float64

```
In [28]: 1 s_future1 = s_future.fillna(0)
```

```
In [29]: 1 s_future1
```

```
Out[29]: 0    0.000000
          1   -4.818263
          2   -4.807703
          3   -4.820040
          4   -4.797442
          ...
          86  -4.917277
          87  -4.890349
          88  -4.931664
          89  -4.974939
          90  -4.998833
          Name: s_current, Length: 91, dtype: float64
```

```
In [30]: 1 s_change = s_future - s_current
          2 s_change = s_change.rename('s_change')
          3 s_change
```

```
Out[30]: 0         NaN
          1   -0.010561
          2    0.012337
          3   -0.022598
          4   -0.010367
          ...
          86  -0.026927
          87    0.041315
          88    0.043275
          89    0.023894
          90   -0.065799
          Name: s_change, Length: 91, dtype: float64
```

```
In [42]: 1 s_change_fitted = np.log(np.abs(s_change)).fillna(0)
```

```
In [43]: 1 s_change_fitted
```

```
Out[43]: 0    0.000000
          1   -4.550623
          2   -4.395167
          3   -3.789915
          4   -4.569119
          ...
          86  -3.614610
          87  -3.186527
          88  -3.140185
          89  -3.734147
          90  -2.721158
          Name: s_change, Length: 91, dtype: float64
```

```
In [31]: 1 df6 = pd.DataFrame(s_change.fillna(0), columns = ['s_change'])
```

```
In [32]: 1 df6
```

```
Out[32]:
```

	s_change
0	0.000000
1	-0.010561
2	0.012337
3	-0.022598
4	-0.010367
...	...
86	-0.026927
87	0.041315
88	0.043275
89	0.023894
90	-0.065799

91 rows × 1 columns

```
In [33]: 1 P1 = len(y['error'])
2 MSE_T = np.sum(np.square(y['error']))/P1
3 MSE_T
```

```
Out[33]: 4.368784909818687e-08
```

```
In [34]: 1 MSE_R = np.sum(np.square([df6['s_change']]))/P1
2 MSE_R
```

```
Out[34]: 0.0005809296335330899
```

```
In [35]: 1 error_R = df6['s_change'].reset_index(drop=True)
2 error_T = y['error'].reset_index(drop=True)
3 tmp = np.square(error_R) - np.square(error_T) - (MSE_R - MSE_T)
4 V_hat = np.sum(np.square(tmp))/P1
5
6 ## Statistic
7 DMW = (MSE_R - MSE_T)/np.sqrt(V_hat/P1)
8
9 print('Since the DMW statitsic is equal to ' + str(DMW) + ',' + ' w
10 print('we reject the null hypothesis that the MP model does not out
```

Since the DMW statitsic is equal to 4.734497242845056, which is more than the critical value (1.28), we reject the null hypothesis that the MP model does not outperform the random walk model.

```
In [44]: 1 tmp2 = np.sum(np.square(s_change_fitted))/P1
2 CW = (MSE_R - MSE_T + tmp2)/np.sqrt(V_hat/P1)
3
4 print('Since the CW statitsic is equal to ' + str(CW) + ',' + ' whi
5 print('we reject the null hypothesis that the MP model does not out
```

Since the CW statitsic is equal to 187938.73511034384, which is more than the critical value (1.28), we reject the null hypothesis that the MP model does not outperform the random walk model.

```
In [ ]: 1
```